

FPGA_and_the_DevCloud

July 1, 2020

1 Walkthrough: FPGA and the DevCloud

This notebook is a demonstration showing you how to request an edge node with an Intel i5 CPU and load a model on the Intel® Arria 10 FPGA using Udacity's workspace integration with Intel's DevCloud. This notebook is just to give you an overview of the process (you won't be writing any code). In the next workspace, you'll be given TODO items to complete.

Below are the six steps we'll walk through in this notebook:

1. Creating a Python script to load the model
2. Creating a job submission script
3. Submitting a job using the qsub command
4. Checking the job status using the liveQStat function
5. Retrieving the output files using the getResults function
6. Viewing the resulting output

Click the **Introduction to FPGA and the DevCloud** button below for a quick overview of the overall process. We'll then walk through each step of the process.

Introduction to FPGA and the DevCloud

IMPORTANT: Set up paths so we can run Dev Cloud utilities You *must* run this every time you enter a Workspace session.

```
In [1]: %env PATH=/opt/conda/bin:/opt/spark-2.4.3-bin-hadoop2.7/bin:/opt/conda/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
import os
import sys
sys.path.insert(0, os.path.abspath('/opt/intel_devcloud_support'))
sys.path.insert(0, os.path.abspath('/opt/intel'))
```

```
env: PATH=/opt/conda/bin:/opt/spark-2.4.3-bin-hadoop2.7/bin:/opt/conda/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

1.1 The Model

We will be using the vehicle-license-plate-detection-barrier-0106 model for this exercise. Remember that to run a model on the Intel® Arria 10 FPGA, we need to use FP16 as the model precision.

The model has already been downloaded for you in the /data/models/intel directory on Intel's DevCloud. We will be using the following filepath during the job submission in **Step 3**:

/data/models/intel/vehicle-license-plate-detection-barrier-0106/FP16/vehicle-
license-plate-detection-barrier-0106

2 Step 1: Creating a Python Script

The first step is to create a Python script that you can use to load the model and perform an inference. I have used the `%%writefile` magic command to create a Python file called `load_model_to_device.py`. This will create a new Python file in the working directory.

Click the **Writing a Python Script** button below for a demonstration.

Writing a Python Script

```
In [2]: %%writefile load_model_to_device.py
```

```
import time
from opencvino.inference_engine import IENetwork
from opencvino.inference_engine import IECore
import argparse

def main(args):
    model=args.model_path
    model_weights=model+'.bin'
    model_structure=model+'.xml'

    start=time.time()
    model=IENetwork(model_structure, model_weights)

    core = IECore()
    net = core.load_network(network=model, device_name=args.device, num_requests=1)
    print(f"Time taken to load model on FPGA = {time.time()-start} seconds")

if __name__=='__main__':
    parser=argparse.ArgumentParser()
    parser.add_argument('--model_path', required=True)
    parser.add_argument('--device', default=None)

    args=parser.parse_args()
    main(args)
```

Overwriting `load_model_to_device.py`

2.1 Step 2: Creating a Job Submission Script

To submit a job to the DevCloud, we need to create a shell script. Similar to the Python script above, I have used the `%%writefile` magic command to create a shell script called `load_fpga_model_job.sh`.

Note: This shell script is the same as the other scripts we've been writing up until this point with one difference. In order to run models on the FPGA, we need to use a bitstream file and program our FPGA. We will do this with the `aocl program` command.

This script does a few things. 1. Writes stdout and stderr to their respective .log files 2. Creates the /output directory 3. Creates DEVICE and MODELPATH variables and assigns their value as the first and second argument passed to the shell script 4. Initialize the environment 5. Load a bitstream file (.aocx) and program the FPGA device. 6. Calls the Python script using the MODELPATH and DEVICE variable values as the command line argument 7. Changes to the /output directory 8. Compresses the stdout.log and stderr.log files to output.tgz

Click the **Creating a Job Submission Script** button below for a demonstration.

Creating a Job Submission Script

```
In [3]: %%writefile load_fpga_model_job.sh
        #!/bin/bash

        exec 1>/output/stdout.log 2>/output/stderr.log

        mkdir -p /output

        DEVICE=$1
        MODELPATH=$2

        export AOCL_BOARD_PACKAGE_ROOT=/opt/intel/opencvino/bitstreams/a10_vision_design_sg2_bits

        source /opt/altera/aocl-pro-rte/aclrte-linux64/init_openc1.sh
        aocl program acl0 /opt/intel/opencvino/bitstreams/a10_vision_design_sg2_bitstreams/2020-2

        export CL_CONTEXT_COMPILER_MODE_INTELFPGA=3

        # Run the load model python script
        python3 load_model_to_device.py --model_path ${MODELPATH} --device ${DEVICE}

        cd /output

        tar zcvf output.tgz stdout.log stderr.log
```

Overwriting load_fpga_model_job.sh

2.2 Step 3: Submitting a Job to Intel's DevCloud

The code below will submit a job to an **IEI Tank-870** edge node with an Intel[®] i5 processor and Intel[®] Arria 10 FPGA. We will load the model on the FPGA.

Note: In order to run inference on the Intel[®] Arria 10 FPGA, FPGA requires the use of the **Heterogenous plugin**, so we'll pass in `HETERO:FPGA,CPU` as the device type argument. We'll learn more about this plugin later in the lesson. As a reminder, when running a model on a FPGA, the model precision we'll need is FP16.

The `!qsub` command takes a few command line arguments: 1. The first argument is the shell script filename - `load_fpga_model_job.sh`. This should always be the first argument. 2. The `-d` flag designates the directory where we want to run our job. We'll be running it in the current directory as denoted by `..`. 3. The `-l` flag designates the node and quantity we want to request. The default quantity is 1, so the 1 after nodes is optional. 4. The `-F` flag let's us pass in a string with all command line arguments we want to pass to our Python script.

Note: There is an optional flag, `-N`, you may see in a few exercises. This is an argument that only works on Intel's DevCloud that allows you to name your job submission. This argument doesn't work in Udacity's workspace integration with Intel's DevCloud.

In the cell below, we assign the returned value of the `!qsub` command to a variable `job_id_core`. This value is an array with a single string.

Once the cell is run, this queues up a job on Intel's DevCloud and prints out the first value of this array below the cell, which is the job id.

Click the **Submitting a Job to Intel's DevCloud** button below for a demonstration.

Submitting a Job to Intel's DevCloud

```
In [4]: job_id_core = !qsub load_fpga_model_job.sh -d . -l nodes=1:tank-870:i5-6500te:iei-mustan
        print(job_id_core[0])
```

JGkAKz71juI7cz0YhRdEHvmVUE0bf9V

2.3 Step 4: Running liveQStat

Running the `liveQStat` function, we can see the live status of our job. Running the this function will lock the cell and poll the job status 10 times. The cell is locked until this finishes polling 10 times or you can interrupt the kernel to stop it by pressing the stop button at the top:



- Q status means our job is currently awaiting an available node
- R status means our job is currently running on the requested node

Note: In the demonstration, it is pointed out that W status means your job is done. This is no longer accurate. Once a job has finished running, it will no longer show in the list when running the `liveQStat` function.

Click the **Running liveQStat** button below for a demonstration.

Running liveQStat

```
In [5]: import liveQStat
        liveQStat.liveQStat()
```

2.4 Step 5: Retrieving Output Files

In this step, we'll be using the `getResults` function to retrieve our job's results. This function takes a few arguments.

1. `job id` - This value is stored in the `job_id_core` variable we created during **Step 3**. Remember that this value is an array with a single string, so we access the string value using `job_id_core[0]`.

2. `filename` - This value should match the filename of the compressed file we have in our `load_fpga_model_job.sh` shell script. In this example, `filename` should be set to `output.tgz`.
3. `blocking` - This is an optional argument and is set to `False` by default. If this is set to `True`, the cell is locked while waiting for the results to come back. There is a status indicator showing the cell is waiting on results.

Note: The `getResults` function is unique to Udacity's workspace integration with Intel's DevCloud. When working on Intel's DevCloud environment, your job's results are automatically retrieved and placed in your working directory.

Click the **Retrieving Output Files** button below for a demonstration.

Retrieving Output Files

```
In [6]: import get_results
```

```
get_results.getResults(job_id_core[0], filename="output.tgz", blocking=True)
```

```
getResults() is blocking until results of the job (id:JGkAKz71juI7cz0YhRdEHvmVUE0bfp9V) are read
Please wait...Success!
```

```
output.tgz was downloaded in the same folder as this notebook.
```

2.5 Step 6: Viewing the Outputs

```
In [7]: !tar xzf output.tgz
```

```
In [8]: !cat stdout.log
```

```
INTELFPGAOCCLSDKROOT is not set
```

```
Using script's current directory (/opt/altera/aocl-pro-rte/aclrte-linux64)
```

```
aoc was not found, but aocl was found. Assuming only RTE is installed.
```

```
AOCL_BOARD_PACKAGE_ROOT is set to /opt/intel/opencvino/bitstreams/a10_vision_design_sg2_bitstream
```

```
Adding /opt/altera/aocl-pro-rte/aclrte-linux64/bin to PATH
```

```
Adding /opt/altera/aocl-pro-rte/aclrte-linux64/linux64/lib to LD_LIBRARY_PATH
```

```
Adding /opt/altera/aocl-pro-rte/aclrte-linux64/host/linux64/lib to LD_LIBRARY_PATH
```

```
Adding /opt/intel/opencvino/bitstreams/a10_vision_design_sg2_bitstreams/BSP/a10_1150_sg2/linux64/
```

```
aocl program: Running program from /opt/intel/opencvino/bitstreams/a10_vision_design_sg2_bitstream
```

```
Programming device: a10gx_2ddr : Intel Vision Accelerator Design with Intel Arria 10 FPGA (acla1
```

```
Program succeed.
```

```
Time taken to load model on FPGA = 3.9904346466064453 seconds
```

```
In [9]: !cat stderr.log
```

```
load_model_to_device.py:13: DeprecationWarning: Reading network using constructor is deprecated.
```

```
model=IENetwork(model_structure, model_weights)
```

```
tar: stdout.log: file changed as we read it
```

```
In [ ]:
```